


Trummer/Analyzing Multimodal Data with Large Language Models



about the author



Immanuel Trummer is an assistant professor for computer science at Cornell University and Head of the Cornell Database Group (<https://www.cs.cornell.edu/research/datasystems>).

Immanuel's papers have been selected for "Best of VLDB", "Best of SIGMOD", for the ACM SIGMOD Research Highlight Award, and for publication in CACM as CACM Research Highlight. Immanuel's online course on data management has reached close to a million views on YouTube. Over the past few years, his group has published extensively on projects that apply large language models in the context of data science. This book is a largely extended version of a series of tutorials, introducing data scientists to large language models, given by Immanuel at premier conferences in data management such as VLDB and BTW.

about the book

Analyzing Multimodal Data with Large Language Models teaches you to use a new generation of AI assistants and Large Language Models (LLMs) to simplify and accelerate common data science tasks. Cornell professor and long-time LLM advocate Immanuel Trummer reveals

techniques he's pioneered for getting the most out of LLMs in data science, including model selection and specialization, techniques for tuning parameters, and reliable prompt templates.

You'll start with an in-depth exploration of how LLMs work. Then, you'll dive into no-code data analysis with LLMs, creating custom operators with the OpenAI Python API, and building complex data analysis pipelines with the cutting edge LangChain framework.

Takeaways:

- Analyze text, tables, images, and audio files
- Extract information from multi-modal data lakes
- Classify, cluster, transform, and query multimodal data
- Build natural language query interfaces over structured data sources
- Use LangChain to build complex data analysis pipelines
- Prompt engineering and model configuration

[minimally qualified reader](#)

For data scientists, data analysts, and others who are interested in making their work easier through the use of artificial intelligence techniques. Readers should have a basic understanding of the Python programming language.

[table of contents](#)

= Part 1: Introduction to Language Models

== Chapter 1 Analyzing data with large language models

=== 1.1 What can language models do?

=== 1.2 Who should read this book

=== 1.3 Why you should read this book

=== 1.4 What you will learn

=== 1.5 How to use language models

=== 1.6 Using language models for data analysis

==== 1.6.1 Using language model directly on data

==== 1.6.2 Data analysis via external tools

=== 1.7 Optimizing model selection

==== 1.7.1 Selecting a model provider

==== 1.7.2 Selecting a specific model

=== 1.8 Optimally using language models

==== 1.8.1 Selecting a model configuration

==== 1.8.2 Choosing a prompt template

=== 1.9 Multi-step data analysis pipelines

=== 1.10 Summary

=== References

== Chapter 2 A chat with ChatGPT

=== 2.1 Accessing the Web interface

=== 2.2 Making Introductions

=== 2.3 Processing text with GPT

=== 2.4 Processing tables with GPT

=== 2.5 Summary

=== References

== Chapter 3 The Transformer

=== 3.1 Tokenization

==== 3.1.1 What is tokenization?

==== 3.1.2 Why learn about tokenization?

==== 3.1.3 Why tokenization is difficult

==== 3.1.4 How GPT tokenizes text

=== 3.2 Attention mechanisms

==== 3.2.1 Queries, keys, and values

==== 3.2.2 A detailed example

==== 3.2.3 What makes attention mechanisms special?

=== 3.3 Encoders

==== 3.3.1 Encoder architecture

==== 3.3.2 Multi-Head Self-Attention

==== 3.3.3 Visualizing attention

==== 3.3.4 Visualizing attention in more detail

=== 3.4 Decoders

==== 3.4.1 Why we need decoders

==== 3.4.2 Decoder architecture

==== 3.4.3 Cross versus self-attention

==== 3.4.4 Turning output into input

==== 3.4.5 Decoder-only models

=== 3.5 Decoding strategies

==== 3.5.1 A simple decoding strategy

==== 3.5.2 Randomizing output via the temperature parameter

==== 3.5.3 Other parameters for randomization

==== 3.5.4 Thinking multiple steps ahead

==== 3.5.5 Beam search

=== 3.6 Summary

=== References

== Chapter 4 Transfer learning

=== 4.1 Transfer learning basics

=== 4.2 Pre-training

==== 4.2.1 Pre-training tasks

==== 4.2.2 Pre-training data

=== 4.3 Specialization

==== 4.3.1 Fine-tuning

==== 4.3.2 Prompting

=== 4.4 Summary

=== References

= Part 2: Data Analysis with Language Models

== Chapter 5 The OpenAI API

=== 5.1 Prerequisites

=== 5.2 Setting up the OpenAI client

=== 5.3 Listing available models

=== 5.4 Making GPT tell us a story

=== 5.5 Customizing text generation

=== 5.6 Chat models

=== 5.7 Summary

== Chapter 6 Analyzing text data

=== 6.1 Preliminaries

=== 6.2 Classification

==== 6.2.1 Overview

==== 6.2.2 Creating prompts

==== 6.2.3 Calling the model

==== 6.2.4 End-to-end classification code

==== 6.2.5 Classifying documents

==== 6.2.6 Running the code

- ==== 6.2.7 Trying out variants
- === 6.3 Text extraction
 - ==== 6.3.1 Overview
 - ==== 6.3.2 Generating prompts
 - ==== 6.3.3 Post-processing
 - ==== 6.3.4 End-to-end extraction code
 - ==== 6.3.5 Trying it out
- === 6.4 Clustering
 - ==== 6.4.1 Overview
 - ==== 6.4.2 Calculating embeddings
 - ==== 6.4.3 Clustering vectors
 - ==== 6.4.4 End-to-end code for text clustering
 - ==== 6.4.5 Trying it out
 - ==== 6.4.6 Other use cases for embedding vectors
- === 6.5 Summary
- === References

- == Chapter 7 Analyzing structured data
 - === 7.1 Overview
 - === 7.2 Setting up an SQLite database
 - === 7.3 The SQL query language
 - === 7.4 A simple question translator
 - ==== 7.4.1 Overview
 - ==== 7.4.2 Code discussion
 - ==== 7.4.3 Trying it out
 - === 7.5 An end-to-end natural language query interface
 - ==== 7.5.1 Overview
 - ==== 7.5.2 Main function
 - ==== 7.5.3 Sub-functions
 - ==== 7.5.4 Trying it out
 - === 7.6 Graph data
 - ==== 7.6.1 Setting up a Neo4j database
 - ==== 7.6.2 The Cypher query language
 - === 7.7 Querying graphs in natural language
 - ==== 7.7.1 Overview
 - ==== 7.7.2 Code discussion

=== 7.8 Summary

== Chapter 8 Analyzing images and videos

=== 8.1 Setup and preliminaries

=== 8.2 Answering questions about images

===== 8.2.1 Specifying multimodal input

===== 8.2.2 Code discussion

===== 8.2.3 Trying it out

=== 8.3 Tagging people in images

===== 8.3.1 Overview

===== 8.3.2 Encoding locally stored images

===== 8.3.3 Sending locally stored images to GPT

===== 8.3.4 The end-to-end implementation

===== 8.3.5 Trying it out

=== 8.4 Generating titles for videos

===== 8.4.1 Overview

===== 8.4.2 Encoding video frames

===== 8.4.3 The end-to-end implementation

===== 8.4.4 Trying it out

=== 8.5 Summary

=== References

= Part 3: Advanced Data Analysis

== Chapter 9 Analyzing audio data

=== 9.1 Preliminaries

=== 9.2 Transcribing audio files

===== 9.2.1 Transcribing speech

===== 9.2.2 End-to-end code

===== 9.2.3 Trying it out

=== 9.3 Querying relational data via voice

===== 9.3.1 Preliminaries

===== 9.3.2 Overview

===== 9.3.3 Recording audio

===== 9.3.4 End-to-end code

===== 9.3.5 Trying it out

=== 9.4 Speech-to-speech translation

==== 9.4.1 Overview

==== 9.4.2 Generating speech

==== 9.4.3 End-to-end code

==== 9.4.4 Trying it out

=== 9.5 Summary

== Chapter 10 Alternative model providers

== Chapter 11 Prompt engineering and model tuning

= Part 4: Conclusion

== Chapter 12 Summary and outlook