Kennedy3/Outlier Detection in Python

about the author

Brett Kennedy is a data scientist with over thirty years' experience in software development and over ten in data science. He has worked in outlier detection related to financial auditing, fraud detection, and social media analysis. He previously led a research team focusing on outlier detection. He lives in Toronto with his spouse and two children.

about the book

This book is an enlightening journey into the world of outlier detection, a crucial yet often overlooked aspect of unsupervised learning. This book showcases how identifying unusual records can help spot errors and other issues, making it a valuable tool in data science projects. It delves into its application in data mining and prediction problems, highlighting its role in identifying data quality issues and checking for data drift. Despite challenges such as lack of ground truth and its subjective nature, the book presents these as intriguing problems, offering a comprehensive understanding of this intellectually stimulating area of data science. It aims to fill the academic gap in outlier detection, making it an essential read for anyone keen to delve into this fascinating field.

minimally qualified reader

Readers would require an introductory background in data science, for example understanding the basics of classification and regression, python, pandas, and scikit-learn.

Target groups:
The book is for people who work with text, time-series, network, or tabular data and wish to identify unusual items within data, or wish to understand how outlier detection works.

table of contents